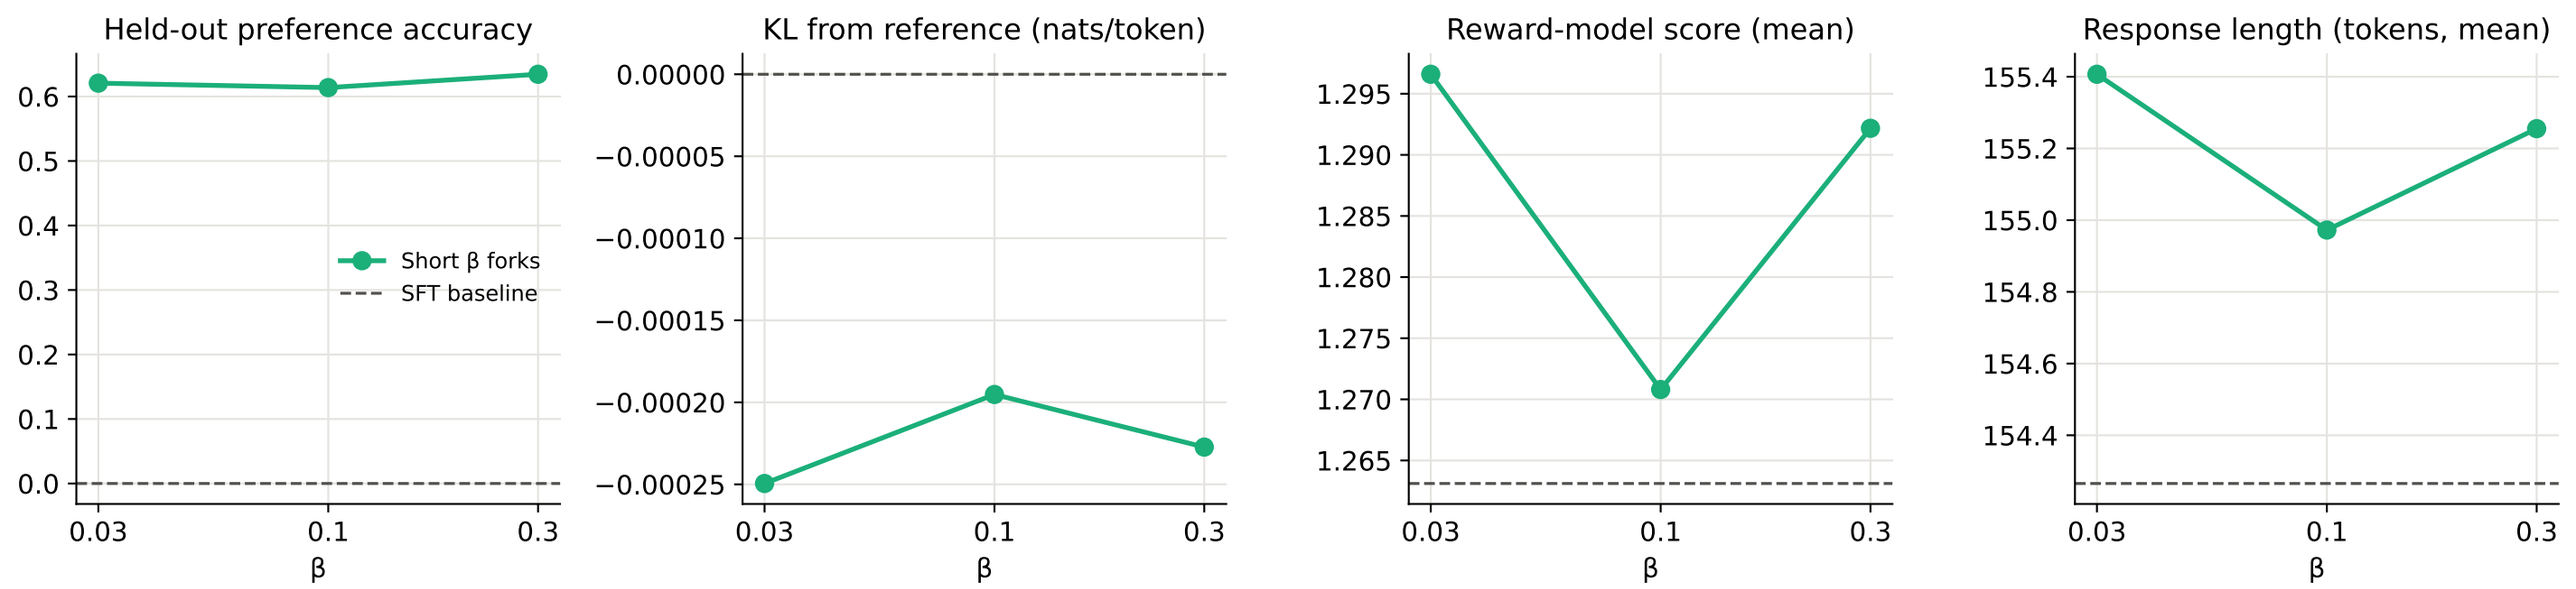
Held-out preference accuracy
0.6
0.5
0.4
0.3
0.2
0.1
0.0
Short β forks
SFT baseline
0.03
0.1
0.3
β
KL from reference (nats/token)
0.00000
−0.00005
−0.00010
−0.00015
−0.00020
−0.00025
0.03
0.1
0.3
β
Reward-model score (mean)
1.295
1.290
1.285
1.280
1.275
1.270
1.265
0.03
0.1
0.3
β
Response length (tokens, mean)
155.4
155.2
155.0
154.8
154.6
154.4
0.03
0.1
0.3
β