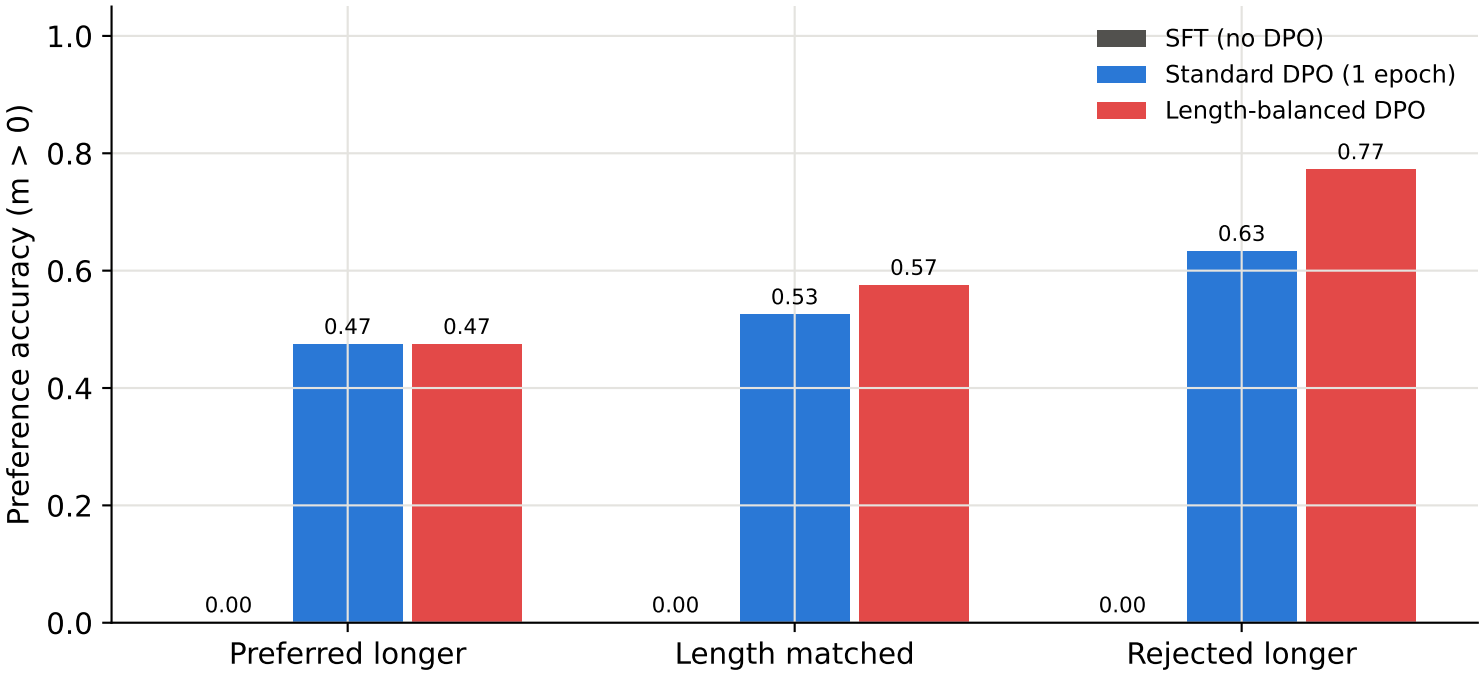
Length-stratified held-out preference accuracy (82 pairs/stratum before filtering)
Preference accuracy (m > 0)
1.0
0.8
0.6
0.4
0.2
0.0
SFT (no DPO)
Standard DPO (1 epoch)
Length-balanced DPO
0.00
0.47
0.47
0.00
0.53
0.57
0.00
0.63
0.77
Preferred longer
Length matched
Rejected longer