Task 1 DPO training curves (standard: 1 epoch; β forks: short budget)
DPO loss (train, per update)
Standard DPO (1 epoch)
β=0.03 (short)
β=0.10 (short)
β=0.30 (short)
Length-balanced DPO
0.80
0.75
0.70
0.65
0
20
40
60
80
Optimizer update
Preference accuracy (train batch)
0.8
0.6
0.4
0.2
0.0
0
20
40
60
80
Optimizer update
Margin m = Δlogratio (nats)
1
0
−1
−2
0
20
40
60
80
Optimizer update
Grad norm (pre-clip)
10
8
6
4
2
0
20
40
60
80
Optimizer update