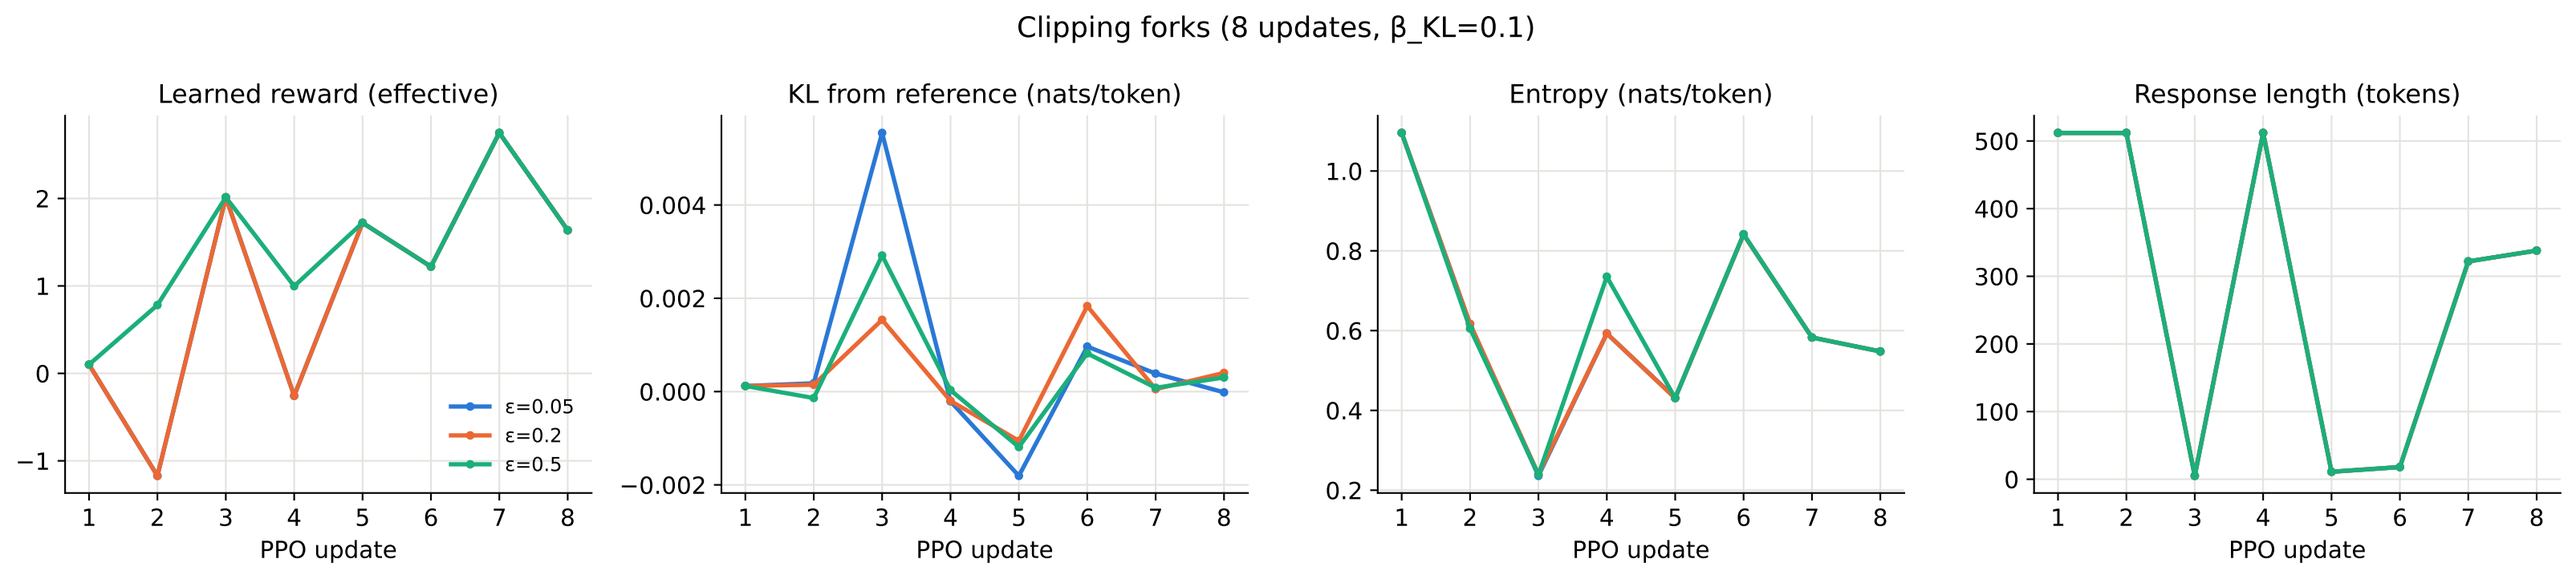
Clipping forks (8 updates, β_KL=0.1)
Learned reward (effective)
KL from reference (nats/token)
Entropy (nats/token)
Response length (tokens)
ε=0.05
ε=0.2
ε=0.5
PPO update
PPO update
PPO update
PPO update