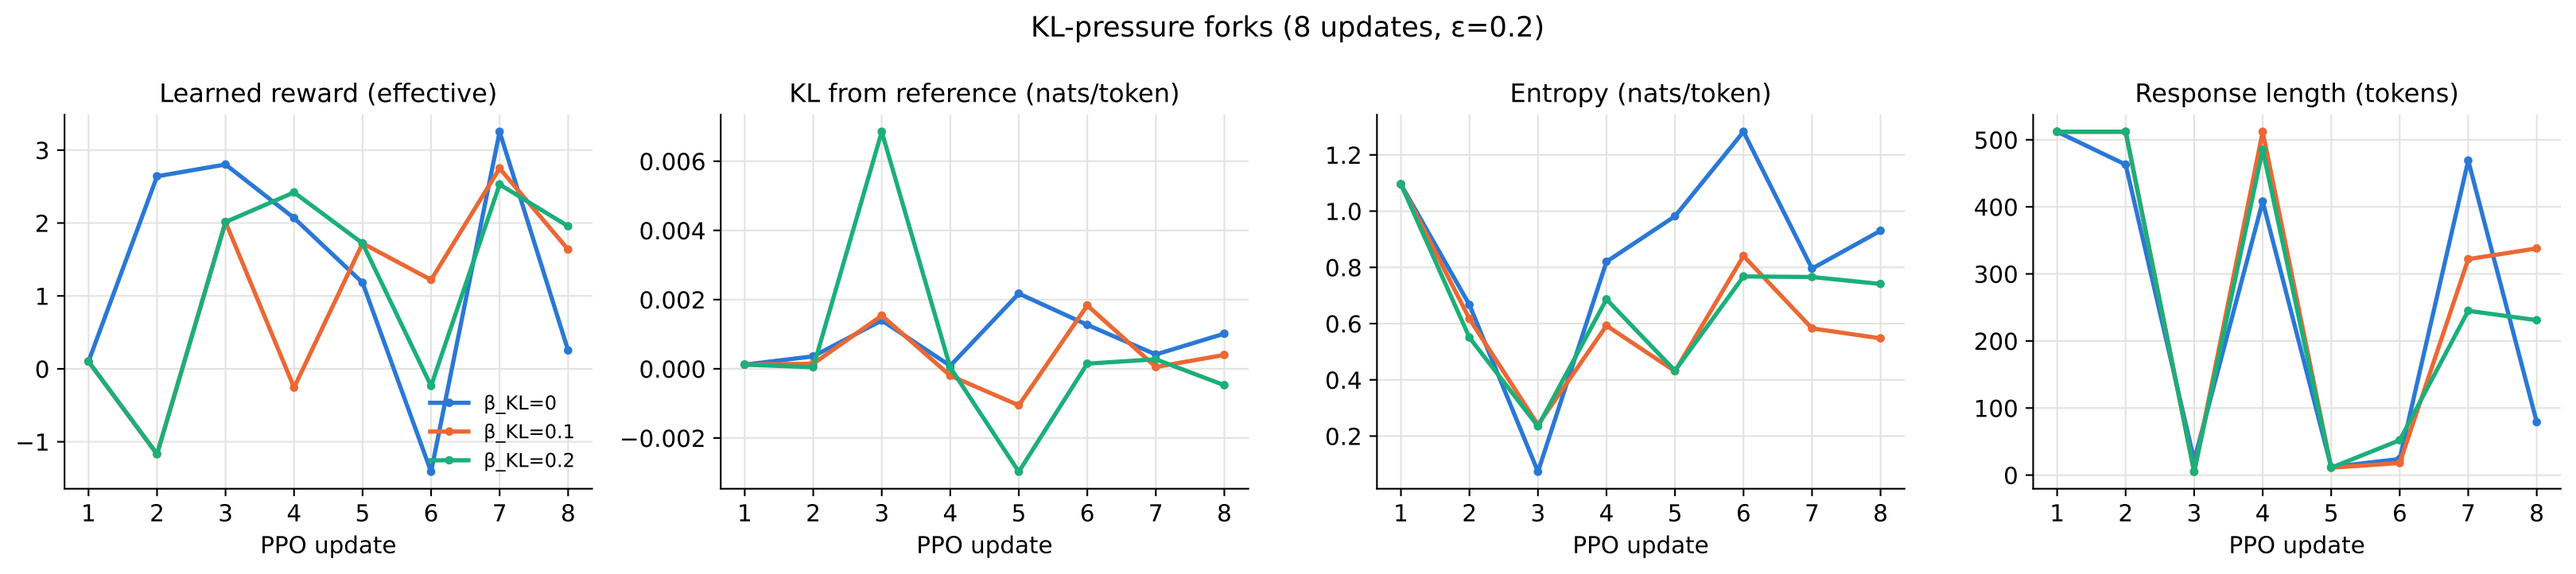
KL-pressure forks (8 updates, ε=0.2)
Learned reward (effective)
KL from reference (nats/token)
Entropy (nats/token)
Response length (tokens)
β_KL=0
β_KL=0.1
β_KL=0.2
PPO update
PPO update
PPO update
PPO update