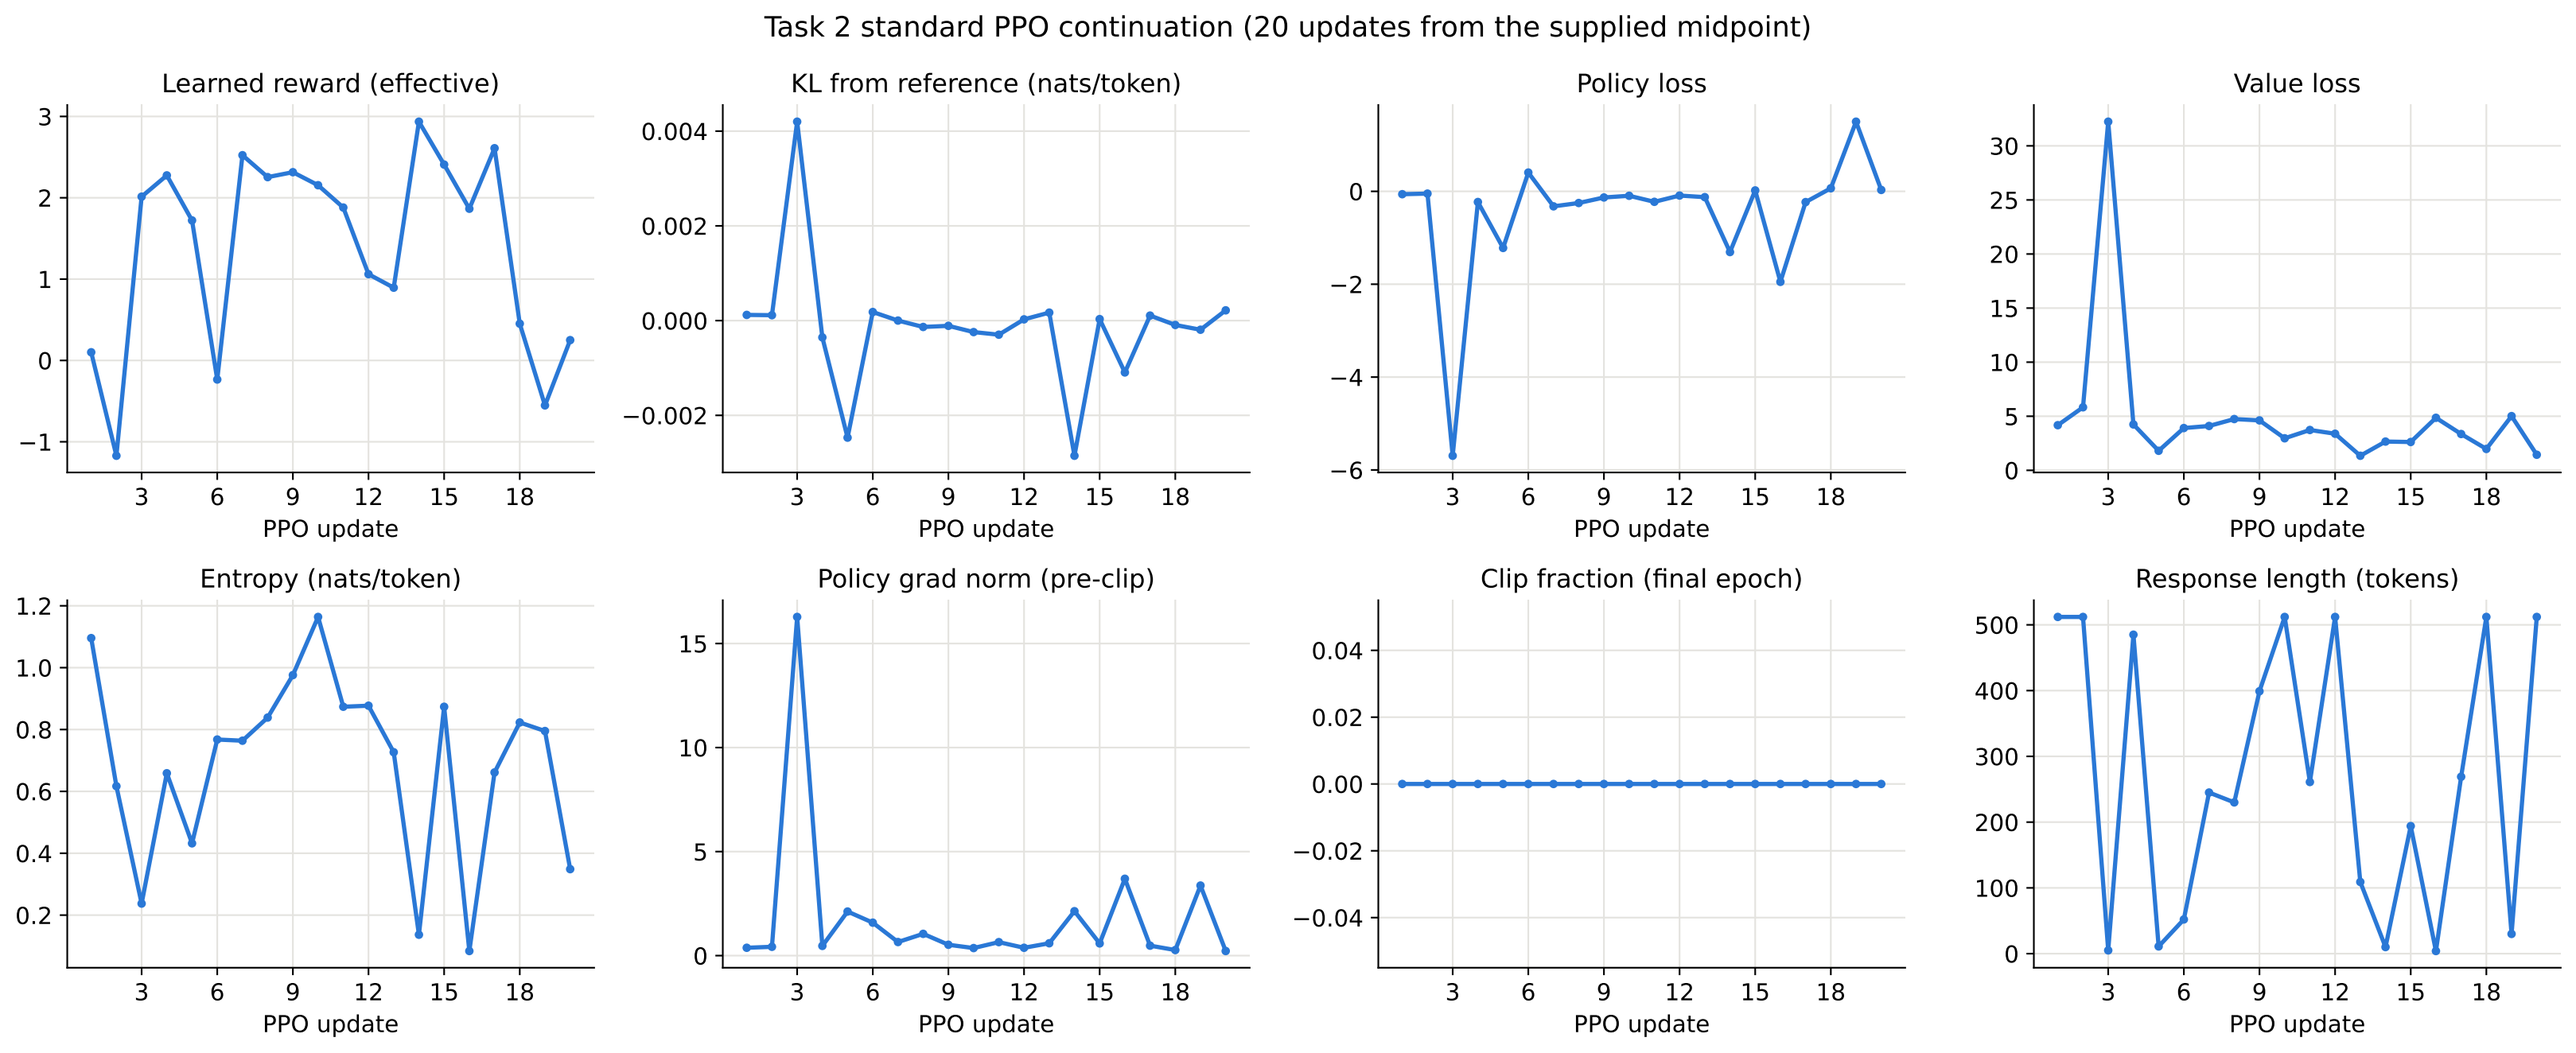
Task 2 standard PPO continuation (20 updates from the supplied midpoint)
Learned reward (effective)
KL from reference (nats/token)
Policy loss
Value loss
Entropy (nats/token)
Policy grad norm (pre-clip)
Clip fraction (final epoch)
Response length (tokens)
PPO update